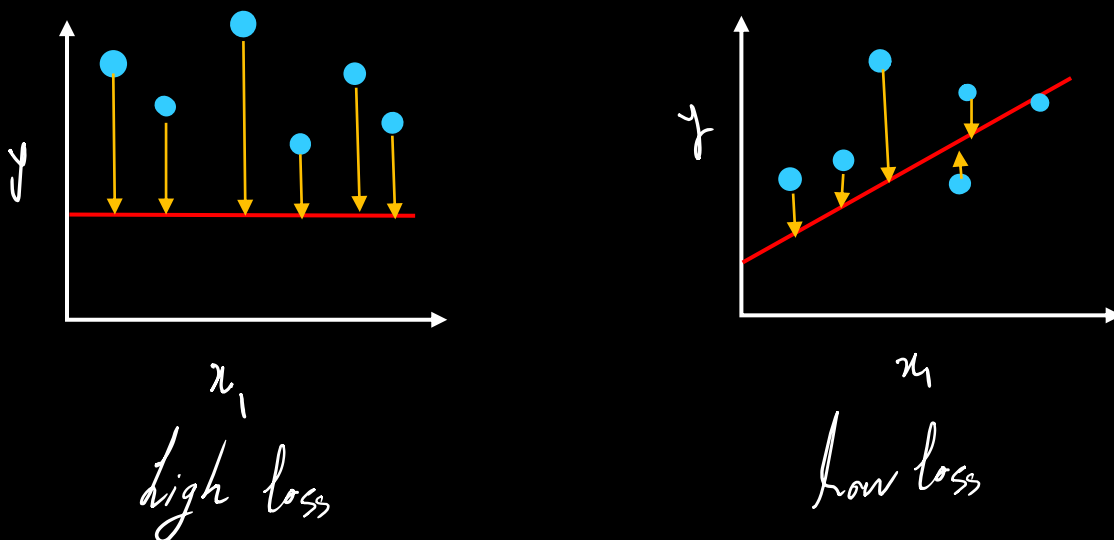

DESCENDING INTO ML: TRAINING AND LOSS

Training a model simply means learning(determining) good values for all the weights and the bias from labeled examples. In supervised learning, a ML algorithm build a model by examining many examples and attempting to find a model that minimizes loss; this process is called empirical risk minimization.

Loss is penalty for a bad prediction. That is, **Loss** is a number indicating how bad the model's prediction was on single example. If model's prediction is perfect, the loss is zero; otherwise, the loss is greater. The goal of training a model is to find a set of weight and biases that have *low* loss, on average, across all examples. For example, Figure below shows a high loss model on the left and a low loss model on the right. Note the following about the figure:

- The arrows represent loss.
- The red line represent prediction.



Notice that the arrows in the left plot are much longer their counterparts in the right plot. Clearly, the line in the right plot is a much better predictive model than the line in the left plot.

You might be wondering whether you could create a mathematical function-a loss function- that would aggregate the individual losses in a meaningful fashion.

Squared loss: a popular loss function

The linear regression models we'll examine here use a loss function called **squared loss** (also known as l_2 loss). The squared loss for a single example is as follows:

$$\begin{aligned} &= \text{the square of difference between the label and the prediction} \\ &= (\text{observation} - \text{prediction}(x))^2 \\ &= (y - y')^2 \end{aligned}$$

Mean square error (MSE) is the average squared loss per example over the whole dataset. To calculate MSE, sum up all the squared losses for individual examples and divide by the number of examples:

$$MSE = \frac{1}{N} \sum_{(x,y) \in D} (y - \text{prediction}(x))^2$$

Where:

- (x, y) is an example in which
 - x is the set of features (for example, chirps/minute, age, gender) that the model uses to make predictions .
 - y is the example's label (for example, temperature).
- $\text{prediction}(x)$ is a function of the weights and bias in combination with the set of features x .
- D is a data set containing many labeled examples, which are (x, y) pairs.
- N is the number of examples in D

Although MSE is commonly-used in ML, it's neither the only practical loss function nor the best loss function for all circumstances.